Attention Is Mostly What You Need

A. Author and B. Writer

LEFTSTART Abstract. Sequence models are usually built from recurrent layers. We propose a simpler design that relies on attention alone. Experiments on two translation tasks show that the model is better in quality while being far easier to train in parallel. Abstract. Sequence models are usually built from recurrent layers. We propose a simpler design that relies on attention alone. Experiments on two translation tasks show that the model is better in quality while being far easier to train in parallel. Abstract. Sequence models are usually built from recurrent layers. We propose a simpler design that relies on attention alone. Experiments on two translation tasks show that the model is better in quality while being far easier to train in parallel. LEFTEND

RIGHTSTART Introduction. Recurrent networks process one position at a time, which makes training slow on long inputs. Attention lets every position look at every other position in one step, so the work can be spread over many processors. Introduction. Recurrent networks process one position at a time, which makes training slow on long inputs. Attention lets every position look at every other position in one step, so the work can be spread over many processors. Introduction. Recurrent networks process one position at a time, which makes training slow on long inputs. Attention lets every position look at every other position in one step, so the work can be spread over many processors. RIGHTEND